



# **HEALTH SCORE PREDICTOR**

**GROUP 06**  
**DATA SCIENCE WITH PYTHON**



OTTO VON GUERICKE  
UNIVERSITÄT  
MAGDEBURG

FAKULTÄT FÜR  
INFORMATIK

- Problem and Motivation
- Data Source
- Notebook map
- Data Exploration
- Feature Engineering
- Balancing and building dataset
- Model
- Results and Insights

# Problem and Motivation

- Limited inspection bandwidth leads to reactive scheduling high risk restaurants are often found late due to fragmented data and manual heuristics.
- Future outcomes (pass/fail/conditional) depend on historical violations, time since last inspection, neighborhood risk, and inspection type hard to weigh consistently without a predictive tool.
- City agencies lack a unified, data-driven way to prioritize inspections, risking missed hotspots and inefficient routes.

- Use SF health inspections plus public signals (Google ratings) to forecast violation risk and proactively target HIGH-risk establishments.
- Provide an inspector-facing dashboard with trends, neighborhood risk, and per-business predictions to inform scheduling and resource allocation.
- Demonstrate ML (RF or XGBoost) can outperform baseline heuristics, improving public health outcomes and fairness with transparent, data-backed decisions.

# Data Source

---

- SF Health Inspections (2020–Present):

Primary labeled records of inspections, violations, and outcomes. These are the raw records which are found separately in the SF health department website [data.sfgov.org/Health-and-Social-Services](http://data.sfgov.org/Health-and-Social-Services)

[!\[\]\(dfbd6b3763a6d1d9afaa974f64e2e4b5\_img.jpg\) HealthInspection\(2020-2023\).csv](#) [!\[\]\(b89ecf30df3dbaee65fa9f1829524a6e\_img.jpg\) HealthInspection\(2024-present\).csv](#)

- Google Places Signals:

Ratings and review counts to proxy hygiene perception. This data is extracted using Google cloud API: Geocoding API and Places API.

[!\[\]\(c694a3ff3b077d76910920a6a1593ab4\_img.jpg\) Google\\_clean.csv](#)

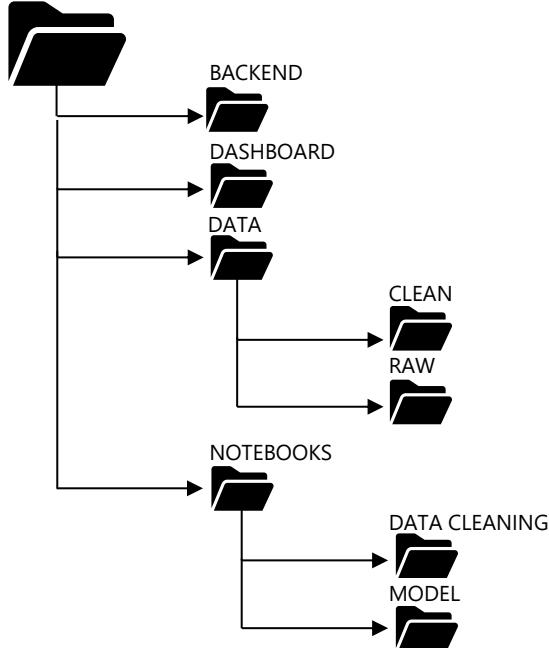
- Neighborhood Crosswalk:

Mapping neighborhoods/ZIPs for spatial features and one-hot encoding. Cleaned up dataset for visualization and model run

[!\[\]\(aa53ad6fea213b8b2226d3077e30533a\_img.jpg\) model\\_dataset.csv](#)

# Notebooks Map

HEALTHSCORE-PREDICTOR



**BACKEND** – Consists of the End points, Data inputs and Model serving

**DASHBOARD** – React frontend framework for visualization and user interaction.

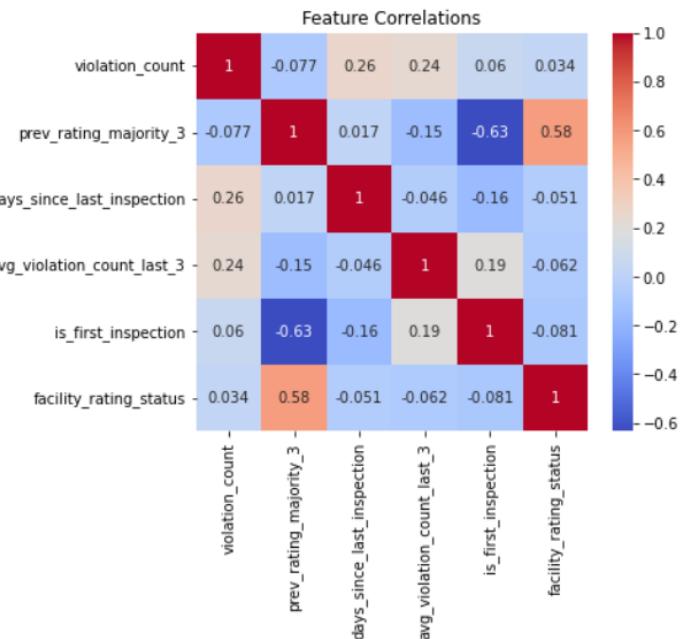
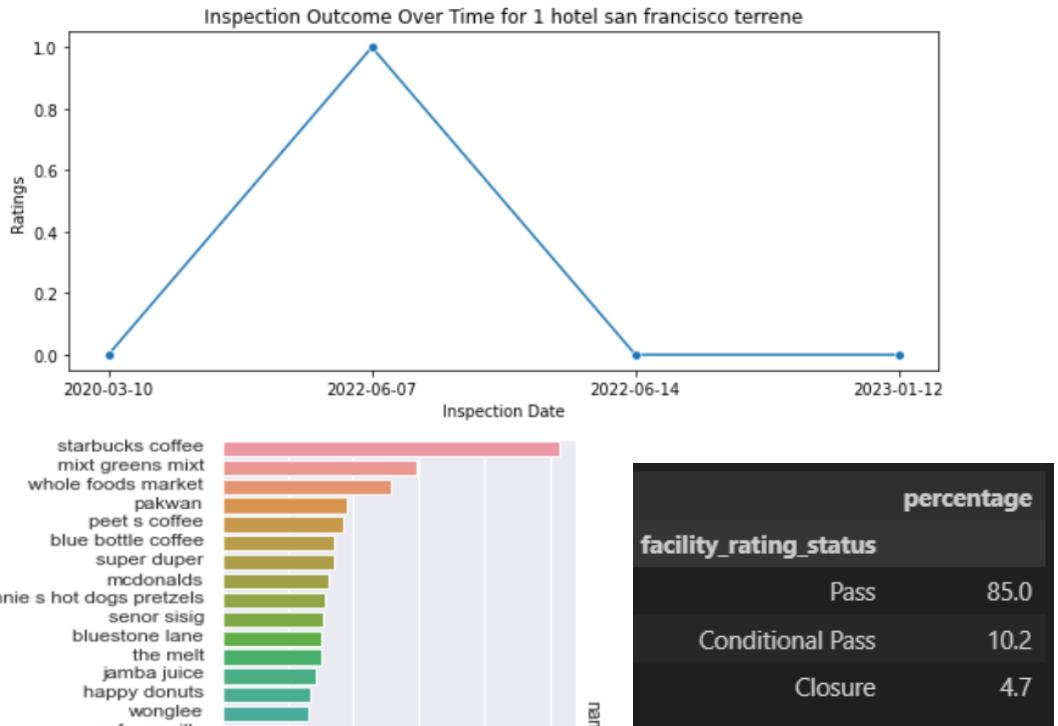
**DATA/CLEAN** – PreProcessed csv file useful for the model run and visual ease

**DATA/RAW** – Untouch dataset extracted.

**NOTEBOOKS/DATACLEANING** – Data cleaning, standardizing, schema-aligned datasets ready for EDA, modeling, and the dashboard.a

**NOTEBOOKS/MODEL** – Execution of our chosen model over the built dataset

# Data Exploration



# Data Cleaning and Feature Engineering

## Text normalization

### **Column names:**

Standardized to lowercase with underscores

### **String trimming:**

Removed leading/trailing spaces in all text columns.

### **Address & name canonicalization:**

Added canonical names dress by normalizing suffixes, removing unit markers, and slugifying text.

### **Corrections:**

Standardized facility status variants (typos in "Conditional Pass").

## Merging & Alignment (2020 vs 2024)

### **Column alignment:**

Renamed 2020 columns to match 2024 schema

### **Column pruning:**

Dropped administrative/geospatial columns not needed for modeling (e.g., inspector, district, point, etc.).

### **Dataset merge:**

Concatenated aligned 2020 and 2024 frames into df\_merge.

## Feature Engineering

### **Recent rating:**

Most common inspection outcome across the last three inspections for the same facility.

### **Recent Violation:**

Average number of violations over the previous three inspections.

### **Inspection gap (in days):**

Elapsed time in days since the facility's previous inspection.

### **First Inspection :** Yes or NO.

## Imputation & Null handling

### **Violations (2024):**

Imputed violation\_count missing to 0 imputed violation\_codes missing to empty string.

### **Violations (merged):**

Filled violation\_count missing to 0 and

created hasViolation\_count to differentiate two datasets

### **Lat and Log impute:**

Data extracted from google\_data and imputed for missing lat and log

# Handling Imbalance

## Balancing Techniques Evaluated

SMOTE

BorderlineSMOTE

Random Over-Sampling

Random Under-Sampling

## Model-Based Imbalance

### Handling

class\_weight used in **Random Forest**

sample\_weight used in **XGBoost**

## Key Observations

- Using class weights on the original dataset outperformed methods that introduced synthetic data
- Random Under-Sampling produced results similar to the original dataset
- Over-sampling and SMOTE degraded model performance Synthetic data introduced noise
- Led to poorer generalization

# Model Benchmarking

## The Challenge: Class Imbalance

The dataset is heavily skewed towards "Pass" ratings. Detecting "Closures" (Class 2) is difficult but critical for public safety.

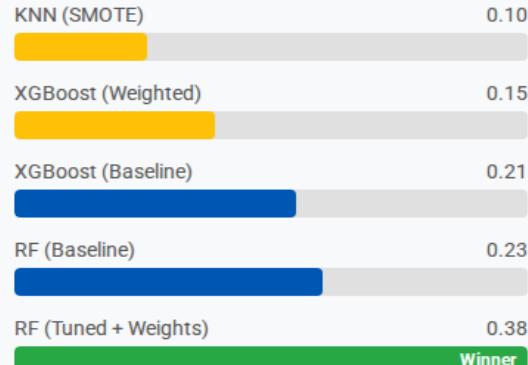
## Selection Rationale

**XGBoost (Baseline):** High accuracy (90%) but failed to detect closures (Recall ~0.25).

**RF (Undersampling):** Improved recall but lost too much training data.

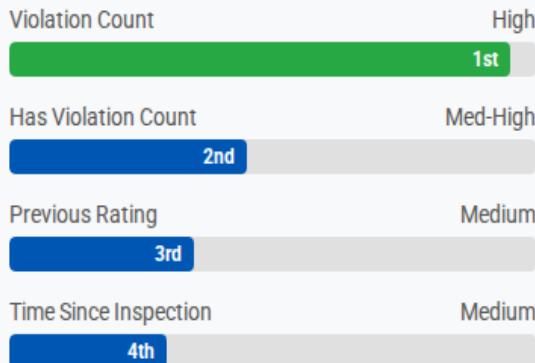
**Winner: RF (Tuned + Weights):** We used custom class weights {0:1, 1:2, 2:4}. This provided the best balance of stability and precision.

Minority Class (Closure) F1-Score



# Results & Strategic Insights

## Key Drivers (Feature Importance)



**94.1%**

OVERALL ACCURACY

**0.85**

CLOSURE PRECISION

**0.68**

OPTIMAL THRESHOLD

## Business Recommendations

**Prioritize Risk:** The sheer volume of current violations is the #1 predictor. Facilities with high violation counts should trigger immediate senior officer review.

**History Matters:** Past performance is predictive. "Conditional Pass" history correlates with future Closures.

**Deployment:** Use the model to flag "At-Risk" facilities *before* inspection to allocate resources efficiently.



# QUESTIONS?